



Practical Work 1 - Design and Optimization of Machine Learning Models

Machine Learning Data
(DAA)

Pg52676 - Catarina Costa

Pg54470 - Fernando Alves

Pg52694 - Marta Aguiar

TABLE OF CONTENTS

1

Dataset Description

2

Data Analysing

3

Data Processing

4

Data Cleaning

5

Data Exploration

6

Modeling

7

Conclusion

DATASET

Customer Personality Analysis

- Customer Persona Understanding
- Targeted Product Modification
- Marketing Strategies Optimizations



DATASET ATTRIBUTES



People

- ID
- Year_Birth
- Education
- Marital_Status
- Income
- Kidhome
- Teenhome
- Dt_Customer
- Recency
- Complain



Products

- MntWines
- MntFruits
- MntMeatProducts
- MntFishProducts
- MntSweetProducts
- MntGoldProds



Promotion

- NumDealsPurchases
- AcceptedComp1
- AcceptedComp2
- AcceptedComp3
- AcceptedComp4
- AcceptedComp5
- Response



Place

- NumWebPurchases
- NumCatalogPurchases
- NumStorePurchases
- NumWebVisitsMonth

DATA ANALYSING

We initially conducted a general analysis of the dataset

data.columns

```
Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome', 'Teenhome', 'Dt_Customer', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts', 'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth', 'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1', 'AcceptedCmp2', 'Complain', 'Z_CostContact', 'Z_Revenue', 'Response'], dtype='object')
```

data.tail()

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth
2235	10870	1967	Graduation	Married	61223.0	0	1	13-06-2013	46	709	...	
2236	4001	1946	PhD	Together	64014.0	2	1	10-06-2014	56	406	...	
2237	7270	1981	Graduation	Divorced	56981.0	0	0	25-01-2014	91	908	...	
2238	8235	1956	Master	Together	69245.0	0	1	24-01-2014	8	428	...	
2239	9405	1954	PhD	Married	52869.0	1	1	15-10-2012	40	84	...	

5 rows × 29 columns

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype  
---  --
0   ID                     2240 non-null  int64  
1   Year_Birth             2240 non-null  int64  
2   Education              2240 non-null  object  
3   Marital_Status         2240 non-null  object  
4   Income                 2216 non-null  float64 
5   Kidhome                2240 non-null  int64  
6   Teenhome               2240 non-null  int64  
7   Dt_Customer            2240 non-null  object  
8   Recency                2240 non-null  int64  
9   MntWines               2240 non-null  int64  
10  MntFruits              2240 non-null  int64  
11  MntMeatProducts        2240 non-null  int64  
12  MntFishProducts        2240 non-null  int64  
13  MntSweetProducts       2240 non-null  int64  
14  MntGoldProds           2240 non-null  int64  
15  NumDealsPurchases      2240 non-null  int64  
16  NumWebPurchases        2240 non-null  int64  
17  NumCatalogPurchases    2240 non-null  int64  
18  NumStorePurchases      2240 non-null  int64  
19  NumWebVisitsMonth      2240 non-null  int64  
20  AcceptedCmp3           2240 non-null  int64  
21  AcceptedCmp4           2240 non-null  int64  
22  AcceptedCmp5           2240 non-null  int64  
23  AcceptedCmp1           2240 non-null  int64  
24  AcceptedCmp2           2240 non-null  int64  
25  Complain               2240 non-null  int64  
26  Z_CostContact           2240 non-null  int64  
27  Z_Revenue              2240 non-null  int64  
28  Response               2240 non-null  int64  
dtypes: float64(1), int64(25), object(3)
memory usage: 507.6+ KB
```

data.head()

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth
0	5524	1957	Graduation	Single	58138.0	0	0	04-09-2012	58	635	...	
1	2174	1954	Graduation	Single	46344.0	1	1	08-03-2014	38	11	...	
2	4141	1965	Graduation	Together	71613.0	0	0	21-08-2013	26	426	...	
3	6182	1984	Graduation	Together	26646.0	1	0	10-02-2014	26	11	...	
4	5324	1981	PhD	Married	58293.0	1	0	19-01-2014	94	173	...	

5 rows × 29 columns

data.shape

data.shape

(2240, 29)

DATA PROCESSING

1) Data conversion to correct dtypes:

- Convert the variable *Dt_Customer* to datetime64[ns] dtype

7	Dt_Customer	2240 non-null	object
---	-------------	---------------	--------



```
data['Dt_Customer'] = pd.to_datetime(data['Dt_Customer'], format='%d-%m-%Y')
```

7	Dt_Customer	2240 non-null	datetime64[ns]
---	-------------	---------------	----------------

2) Feature Engineering:

- Creation of the variables:
 - *Age*: Costumer's age in 2023
 - *Kids*: Sum of the Teenhome and Kidhome variables -> Represents the total number of children
 - *Spent*: Total spending on various items
 - *Is_Parent*: Binary variable

DATA PROCESSING

5) Drop data:

- We decided to drop some unused data for better organization:

```
to_drop = ["Year_Birth", "Z_CostContact", "Z_Revenue",  
           "Kidhome", "Teenhome", "AcceptedCmp1",  
           "AcceptedCmp2", "AcceptedCmp3", "AcceptedCmp4",  
           "AcceptedCmp5", "Complain", "Response"]
```


DATA CLEANING

Age:

`data.describe()`

`data["Age"].value_counts()`

`data_clean = data[data["Age"] <= 100]`

p4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	Response	Age	Kids	Is_Parent	Spent
00	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000
54	0.072768	0.064286	0.013393	0.009375	0.149107	54.194196	0.950446	0.715179	605.798214
00	0.000000	0.000000	0.000000	0.000000	0.000000	27.000000	0.000000	0.000000	5.000000
00	0.000000	0.000000	0.000000	0.000000	0.000000	46.000000	0.000000	0.000000	68.750000
00	0.000000	0.000000	0.000000	0.000000	0.000000	53.000000	1.000000	1.000000	396.000000
00	0.000000	0.000000	0.000000	0.000000	0.000000	64.000000	1.000000	1.000000	1045.500000
00	1.000000	1.000000	1.000000	1.000000	1.000000	130.000000	3.000000	1.000000	2525.000000
28	0.259813	0.245316	0.114976	0.096391	0.356274	11.984069	0.751803	0.451430	602.249288

```
124    1
82     1
130    1
123    1
83     1
Name: count, dtype: int64
```

Since there is a total of three people with the age above 100 years, we decided to drop all the data that had a "Age" bigger than 100.

```
data_clean = data.drop(data.loc[data['Age'] > 100].index, inplace = True)
print(data_clean)
```

DATA CLEANING

data.nunique

```
data.nunique()
```

ID	2237
Education	3
Marital_Status	2
Income	1971
Dt_Customer	663
Recency	100
Wines	775
Fruits	158
Meat	557
Fish	182
Sweets	177
Gold	213
DealsPurch	15
WebPurch	15
CatalogPurch	14
StorePurch	14
WebVisits	16
Age	56
Kids	4
Is_Parent	2
Spent	1054
dtype:	int64

data.isna().null()

```
: data.isna().any()
```

: ID	False
Education	False
Marital_Status	False
Income	True
Dt_Customer	False
Recency	False
Wines	False
Fruits	False
Meat	False
Fish	False
Sweets	False
Gold	False
DealsPurch	False
WebPurch	False
CatalogPurch	False
StorePurch	False
WebVisits	False
Age	False
Kids	False
Is_Parent	False
Spent	False
dtype:	bool

data.isnull().sum()

```
print("Total de valores nulos ")  
print(data.isnull().sum())
```

Total de valores nulos	
ID	0
Education	0
Marital_Status	0
Income	24
Dt_Customer	0
Recency	0
Wines	0
Fruits	0
Meat	0
Fish	0
Sweets	0
Gold	0
DealsPurch	0
WebPurch	0
CatalogPurch	0
StorePurch	0
WebVisits	0
Age	0
Kids	0
Is_Parent	0
Spent	0
dtype:	int64

data.duplicated().sum()

Check if the dataset has duplicate values

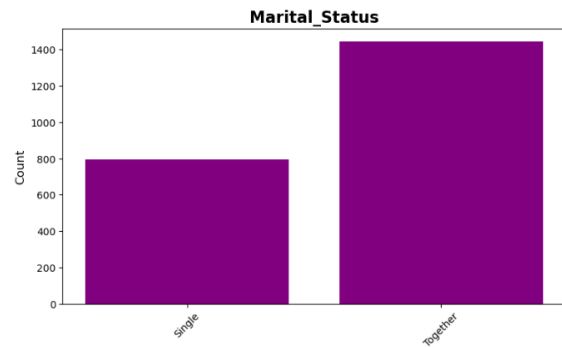

```
data.duplicated().sum()
```

0

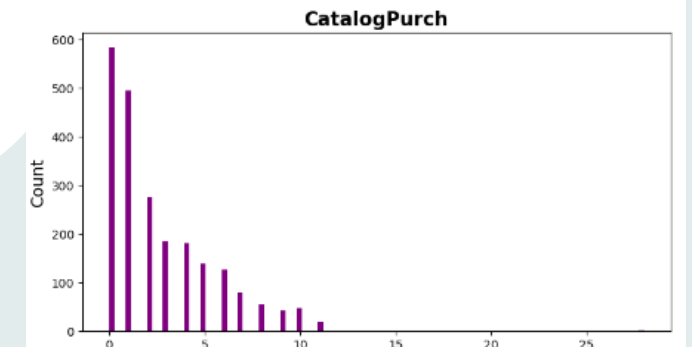









DATA EXPLORATION

Visual Analysis: Distribution

Categorical

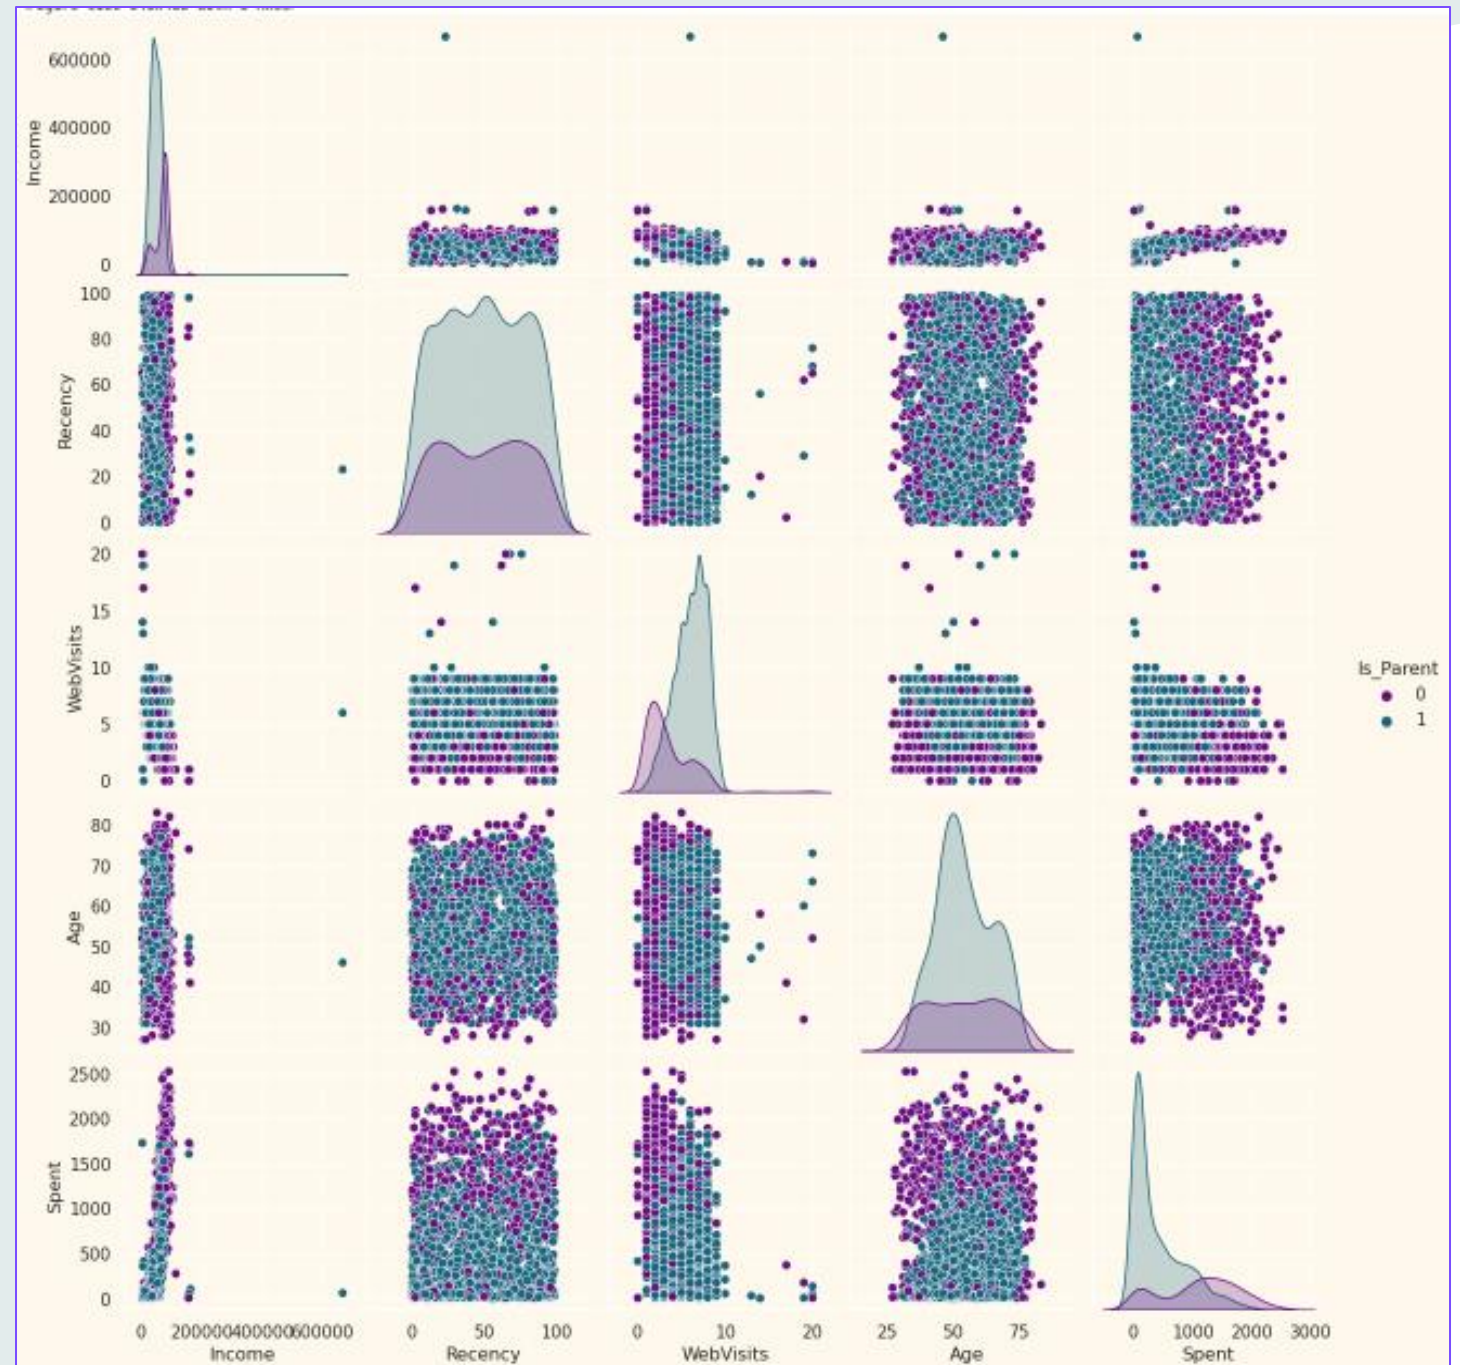


Numerical



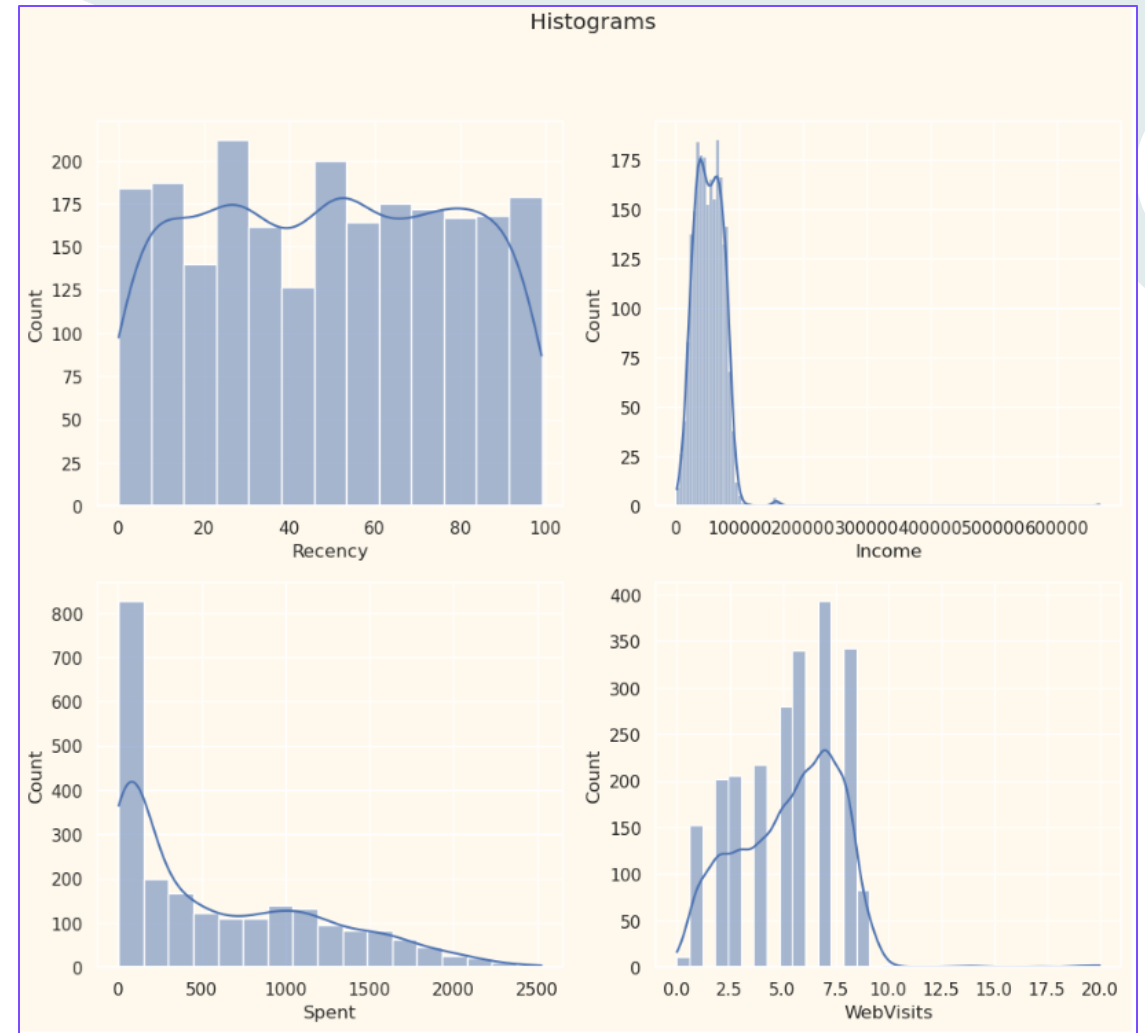
DATA EXPLORATION

Visual Analysis: Multivariable
Analysis







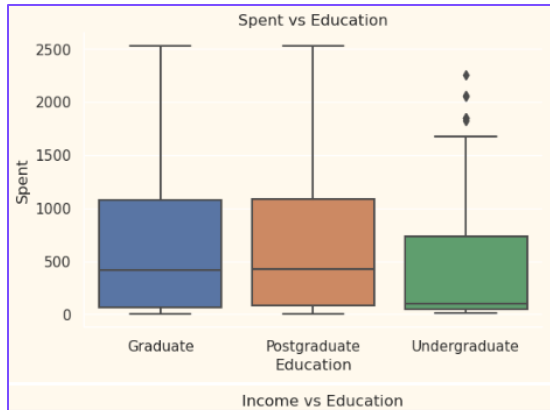
DATA EXPLORATION

Dispersion Graphs:



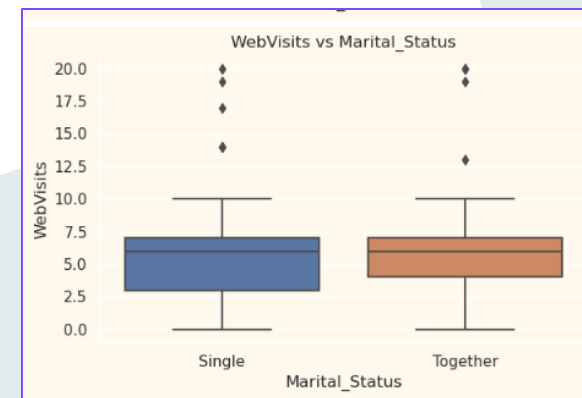
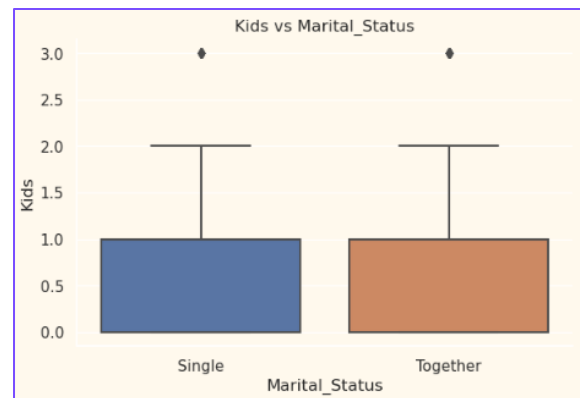
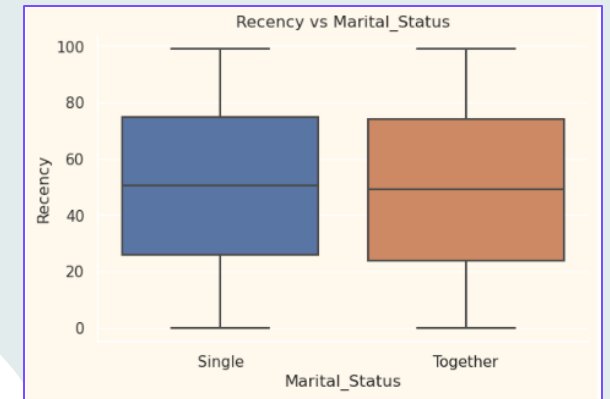


DATA EXPLORATION

Numerical vs Categorical variables: *Education*




DATA EXPLORATION

Numerical vs Categorical variables: *Marital_Status*



DATA EXPLORATION

Correlation Matrix




MODELING – CHOSEN MODELS

A pink rounded rectangle with a teal border, containing the text 'Support Vector Classifier'.

Support
Vector
Classifier

An orange rounded rectangle with a teal border, containing the text 'Grid Search'.

Grid Search

A teal rounded rectangle with a teal border, containing the text 'Linear Regression'.

Linear
Regression

A purple rounded rectangle with a teal border, containing the text 'Decision Trees'.

Decision
Trees

MODELING – SVC and GRIDSEARCH

Regarding this two regression models, the data was treated identically.

1) Data Treatment:

- We created a new dataframe *df_feat* by removing the categorical data and the following target columns, for further training.

```
df_feat = data.drop(['Education', 'Dt_Customer', 'Marital_Status'], axis = 1)
df_target = data['Marital_Status']
```

2) Split Data Test:

- We split the features ('X') and target variable ('y') into training and testing sets;
- 25% of the data is reserved for testing;
- Using *random_state=2022* ensures consistent results across different runs.

```
X_train, X_test, y_train, y_test = train_test_split(df_feat, df_target, test_size=0.25, random_state=2022)
```

MODELING – SUPPORT VECTOR CLASSIFIER (SVC)

1) Data Training:

- Using Cross Validation approach with 10 folds

```
cross_valid_model = SVC(random_state=2022)
scores = cross_val_score(cross_valid_model, df_feat, df_target, cv=10)
```

0.65 accuracy with a standard deviation of 0.00

- Without Cross Validation

```
svc_model = SVC(random_state=2022, class_weight='balanced')
```

```
svc_model.fit(X_train, y_train)
```

2) Data Predictions:

```
svc_predictions = svc_model.predict(X_test)
```

```
print("%0.2f accuracy" % (accuracy_score(y_test, svc_predictions)))
```

0.51 accuracy

MODELING – GRIDSEARCH

1) Data Training:

- We used *RandomForestTreeClassifier*
- GridSearch with Random Forest Classifier involves using GridSearchCV, a hyperparameter tuning technique, to find the best hyperparameters for the Random Forest Classifier model. Random Forest builds multiple decision trees during training and outputs the mode of the classes for predictions.

```
gs_model = RandomForestClassifier(class_weight='balanced', random_state=2023)
```

2) Data Predictions:


```
gs_predictions = gs_model.predict(X_test)
```

```
gs_model.fit(X_train, y_train)
```

MODELING – SVC vs GRIDSEARCH


Support Vector Classification

Classification Report:				
	precision	recall	f1-score	support
0	0.29	0.28	0.28	193
1	0.62	0.63	0.63	361
accuracy			0.51	554
macro avg	0.46	0.46	0.46	554
weighted avg	0.51	0.51	0.51	554



GridSearch

Classification Report:				
	precision	recall	f1-score	support
0	0.62	0.15	0.25	202
1	0.66	0.95	0.78	352
accuracy			0.66	554
macro avg	0.64	0.55	0.51	554
weighted avg	0.65	0.66	0.58	554



Conclusion:

- Classification Report**

- The model GridSearch has better values of precision and recall overall.

- Confusion Matrix**

- The percentage of true values (true positives and true negatives) predicted on the GridSearch model (~65.7%) are higher than the SVC model percentage (~51%).

We can conclude that the model GridSearch was a better approach to train and test data than the SVC.

MODELING – LINEAR REGRESSION

1) Data Treatment:

- We defined a feature matrix X by dropping the following columns;
- Our target variable, y , is set to *Spent*.

```
X = data.drop(['Education', 'Dt_Customer', 'Marital_Status'] + ['Kidhome', 'Teenhome'] + ['Wines', 'Fruits', 'Meat', 'Fish', 'Sweets', 'Gold'] + ['Spent'], axis=1)
y = data['Spent']
```

2) Split Data Test:

- We split the features ('X') and target variable ('y') into training and testing sets;
- 25% of the data is reserved for testing;
- Using *Random_state=2023* ensures consistent results across different runs.

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=2023)
```

MODELING – LINEAR REGRESSION

3) Model Evaluation:

Initialization of a Linear Regression model object

```
lm = LinearRegression()  
lm.fit(X_train,y_train)
```

Initialization of a Linear Regression model object

```
print(lm.intercept_)  
-59.68825869202374
```

Model coefficients

```
coeff_df = pd.DataFrame(lm.coef_,X.columns,columns=['Coefficient'])  
coeff_df
```


	Coefficient
Income	0.004122
Recency	0.100997
DealsPurch	-8.994111
WebPurch	37.424931
CatalogPurch	85.003446
StorePurch	42.733041
WebVisits	-2.426086
Age	-0.494622
Kids	-129.065452

MODELING – LINEAR REGRESSION

4) Model Predictions:

```
predictions = lm.predict(X_test)
```

```
Mean Absolute Error: 187.23436989640257  
Mean Squared Error: 75244.9668507371  
RMSE: 274.3081603794118
```



After exploring the data using the linear regression model, we can conclude that this model is not the most suitable for exploring the data in the selected dataset. For this reason, we decided to continue exploring the data with other models.

MODELING – LINEAR REGRESSION

3) Model Evaluation: (Continuation)

Interpreting the coefficients:

- *Age*: Shows moderate negative impact.
- *WebVisits*: Negatively influences, but to a lesser extent.
- *Kids*: Displays a substantial negative impact having the largest negative coefficient.
- *DealsPurch*: Indicates a significant negative impact.

DealsPurch	-8.994111
WebVisits	-2.426086
Age	-0.494622
Kids	-129.065452

After observing the impact of the '**Age**', '**Kids**', '**DealsPurch**', '**WebVisits**' columns, we opted to remove them, since their impact on the model's predictive capacity appeared limited.

```
X = X.drop(['Age', 'Kids', 'DealsPurch', 'WebVisits'], axis=1)
```

Next, we splitted the data for testing again.

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=2023)
```

MODELING – DECISION TREES

1) Data Treatment:

- We created a new dataframe *data2* by copying the original dataset *data*.
- We defined a feature matrix *X* by dropping columns and setting the target variable *y* to *Spent*.

```
data2 = data.copy()
X = data2.drop(['Marital_Status', 'Education', 'Dt_Customer', 'Spent'], axis=1)
y = data['Spent'].to_frame()
```

2) Split Data Test:

- We split the features ('X') and target variable ('y') into training and testing sets;
- 25% of the data is reserved for testing,
- Using *random_state=2024* ensures consistent results across different runs.

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=2024)
```

MODELING – DECISION TREES

3) Model Analysis and Evaluation:

```
import time

def analyse_model(model, model_name, y_train_analise=y_train):
    start_time = time.time()
    model.fit(X_train, y_train_analise)
    predictions = model.predict(X_test)
    print("time - {}".format(time.time()-start_time))
    predictions = predictions.reshape(len(predictions),1 )
    # Métricas
    print(model)
    #print("Parâmetros:")
    #print(model.get_params())
    print("Mean Absolute Error: ", mean_absolute_error(y_test, predictions))
    print("Mean Squared Error: ", mean_squared_error(y_test, predictions, squared=True))
    print('RMSE: ', np.sqrt(metrics.mean_squared_error(y_test, predictions)))

    r2 = r2_score(y_test, predictions)
    print('R² Score: ', r2)

# sns.displot(y_test-predictions)
# plt.show()
ax = plt.axes()
ax.plot([0, 500, 1000, 2000, 2500, 3000], [0, 500, 1000, 2000, 2500, 3000], 'r')
plt.scatter(y_test, predictions)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title(model_name)
plt.annotate(f'R² Score: {r2:.4f}', xy=(0.7, 0.1), xycoords='axes fraction', fontsize=10, ha='center', color='blue')
plt.show()
```

Interpretation:

Usage - The *analyse_model* function can be used to analyse and visualize the performance of different regression models.

Time Execution - used to measure the execution time of fitting the model and making predictions.

Regression Metrics - Mean Absolute Error, Mean Squared Error, and RMSE (Root Mean Squared Error) are calculated and printed.

R² Score - The R² score measures the goodness of fit.

MODELING – DECISION TREES


4) Model Selection and Evaluation:

```
dtr = DecisionTreeRegressor(random_state=2024)
rfr = RandomForestRegressor(n_estimators=20, max_depth=10, criterion='squared_error', random_state=2024)
gbr = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=2024)

analyse_model(dtr, "Decision Tree Regressor")
analyse_model(rfr, "Random Forest Regressor", y_train.values.ravel())
analyse_model(gbr, "Gradient Boosting Regressor")
```

From the three models (Decision Tree Regressor, Random Forest Regressor, Gradient Boosting Regressor), the Gradient Boosting outperforms in terms of predictive accuracy (lowest *MAE*, *MSE*, *RMSE*, and highest R^2).

The choice of the best model depends on the trade-off between computational efficiency and predictive performance. Considering the balance of performance and execution time, Gradient Boosting seems to be a strong candidate for this regression task.



```
time - 0.6810247898101807
GradientBoostingRegressor(random_state=2024)
Mean Absolute Error: 26.949907129257323
Mean Squared Error: 1583.0128107861642
RMSE: 39.78709352021286
R² Score: 0.99552490255378
```

CONCLUSION

- After preparing the data analysis, processing, cleaning and exploration, we were able to have a global view of how the data in the chosen dataset is organized;
- We were able to see which factors influence some people to spend more or less;
- With the training and testing models, we can say that the **DecisionTrees** model was the model that made the best data prediction.

REFERENCES

Dataset: <https://www.kaggle.com/datasets/imakash3011/customer-personality-analysis>